

# Implementation Manual – Conflict in Ethiopia Before and After Abiy Ahmed

## 0. Purpose & Overview

This manual tells Cursor AI **exactly how to build a full analysis pipeline** for the blog post:

**"Conflict in Ethiopia: Before and After Abiy Ahmed"**

We will:

- Use **ACLED** as the primary quantitative dataset.
- Define the **cut-off date** as the day Abiy Ahmed was sworn in as Prime Minister: **2018-04-02** (configurable).
- Build a **reproducible Python analysis stack** (data ingestion → cleaning → feature engineering → spatio-temporal analysis → figures & tables → blog-ready outputs).
- Generate **publication-quality figures and tables** suitable for EthiopiaConflictAnalytics.org.

**Instruction to Cursor AI (global mindset):** - Always write **clean, modular, well-documented Python code**. - Prefer **functions and small modules** over giant scripts. - Add **docstrings, type hints, and basic tests** where helpful. - Use **sensible defaults**, but make the analysis configurable via a central `config.py`.

---

## 1. Project Structure & Environment

### 1.1 Create project folder & structure

**Goal:** Set up a clean repo where all code, data, and outputs are organized.

Suggested structure:

```
ethiopia_conflict_pre_post_abiy/
├── data/
│   ├── raw/
│   ├── interim/
│   └── processed/
├── geo/
│   └── ethiopia_admin_boundaries/    # shapefiles / GeoJSON
├── notebooks/
└── src/
    └── __init__.py
```

```

|   ┌── config.py
|   ┌── acled_client.py
|   ┌── data_prep.py
|   ┌── features.py
|   ┌── analysis_time_series.py
|   ┌── analysis_spatiotemporal.py
|   ┌── analysis_statistical.py
|   ┌── visualization.py
|   └── utils_logging.py
└── reports/
    ├── figures/
    └── tables/
├── .env
└── pyproject.toml or requirements.txt
└── README.md

```

### Atomic Prompt 1.1 – Create project skeleton

*"You are an assistant helping me set up a Python project for conflict analysis using ACLED data. Create a folder structure (you can describe it in text and generate any starter files) with the following layout: `data/raw`, `data/interim`, `data/processed`, `geo/ethiopia_admin_boundaries`, `notebooks/`, `src/` with modules: `config.py`, `acled_client.py`, `data_prep.py`, `features.py`, `analysis_time_series.py`, `analysis_spatiotemporal.py`, `analysis_statistical.py`, `visualization.py`, `utils_logging.py`, plus `reports/figures`, `reports/tables`, `.env`, `README.md`, and a dependency file (`pyproject.toml` or `requirements.txt`). For each Python module, add a minimal docstring explaining its purpose in the pre/post-Abiy conflict analysis pipeline."*

## 1.2 Set up Python environment & dependencies

Recommended stack:

- `python` (3.10+)
- Core: `pandas`, `numpy`, `python-dotenv`, `requests` (or `httpx`), `tqdm`
- Geo: `geopandas`, `shapely`, `pyproj`
- Plots: `matplotlib`, `seaborn` (for convenience)
- Stats: `statsmodels`, `scipy`

### Atomic Prompt 1.2 – Define dependencies

*"Generate a `requirements.txt` (or `pyproject.toml`) section for this project. It should include: `pandas`, `numpy`, `python-dotenv`, `requests`, `tqdm`, `matplotlib`, `seaborn`, `geopandas`, `shapely`, `pyproj`, `statsmodels`, `scipy`. Also update the `README.md` with step-by-step instructions for creating a virtual environment and installing these dependencies, assuming the project is called `ethiopia_conflict_pre_post_abiy`."*

---

## 2. Configuration & Secrets

### 2.1 Central configuration file

**Goal:** Make the analysis configurable from one place.

Add `src/config.py` with variables like:

```
from datetime import date
from pathlib import Path

PROJECT_ROOT = Path(__file__).resolve().parents[1]

AB_IY_CUTOFF_DATE = date(2018, 4, 2)
COUNTRY = "Ethiopia"

ACLED_BASE_URL = "https://api.acleddata.com/acled/read"
DATA_DIR = PROJECT_ROOT / "data"
GEO_DIR = PROJECT_ROOT / "geo"
REPORTS_DIR = PROJECT_ROOT / "reports"

START_YEAR = AB_IY_CUTOFF_DATE.year - 5
# five calendar years before the cutoff (e.g., 2013)
END_YEAR = 2025      # configurable
END_YEAR = 2025      # configurable
```

(**Note:** fix any whitespace issue above when coding.)

**Atomic Prompt 2.1 – Implement `config.py`**

*"Create a `src/config.py` module that centralizes configuration for the Ethiopia conflict pre-post Abiy analysis. Include: `PROJECT_ROOT`, `DATA_DIR`, `GEO_DIR`, `REPORTS_DIR`, `ACLED_BASE_URL`, `COUNTRY`, `AB_IY_CUTOFF_DATE` (2018-04-02 as a Python `date`), and dynamic `START_YEAR`/`END_YEAR` parameters for ACLED pulls. By default, set `START_YEAR = AB_IY_CUTOFF_DATE.year - 5` so the analysis window begins five calendar years before the cutoff date (e.g., 2013 for a 2018 cutoff), and choose `END_YEAR` as the latest year with stable ACLED coverage. Use `pathlib.Path` and add type hints and docstrings so other modules can import from `config` cleanly."*

### 2.2 Environment variables and secrets

Use a `.env` file to store ACLED credentials:

```
ACLED_EMAIL=your_registered_email@example.org  
ACLED_API_KEY=your_acled_api_key_here
```

#### Atomic Prompt 2.2 – Load secrets via dotenv

"In `src/config.py`, integrate `python-dotenv` to load environment variables from a `.env` file in the project root. Add helper functions or constants: `ACLED_EMAIL`, `ACLED_API_KEY`. If they are missing, raise a clear error message guiding the user to create the `.env` file. Do not print secrets to logs."

### 3. Connecting to the ACLED API

#### 3.1 Read the ACLED API documentation (conceptual)

Cursor doesn't need to fetch the docs, but code must follow the pattern:

- Base endpoint: `https://api.acleddata.com/acled/read`
- Important query parameters: `key`, `email`, `country`, `year`, `limit`, `page`, potentially `event_date`, `region`, etc.
- Responses are paginated.

#### 3.2 Build `acled_client.py`

##### Responsibilities:

- Provide functions to **download ACLED events for Ethiopia** for a given set of years.
- Handle **pagination** and **rate-limiting** politely.
- Cache raw responses as CSV in `data/raw/` for reproducibility.

Suggested functions:

```
fetch_acled_page(params: dict) -> dict  
fetch_acled_for_year(year: int) -> pd.DataFrame  
fetch_acled_range(start_year: int, end_year: int) -> pd.DataFrame  
save_raw(df, path)
```

#### Atomic Prompt 3.2 – Implement ACLED client

"In `src/acled_client.py`, implement a robust ACLED API client. Requirements: - Use `requests` to call `config.ACLED_BASE_URL` with `key` and `email` from `config`. - Implement `fetch_acled_for_year(year: int) -> pd.DataFrame` that retrieves all ACLED records for Ethiopia for that year, handling pagination (e.g., using `page` and `limit` parameters). - Implement `fetch_acled_range(start_year: int, end_year: int) ->`

`pd.DataFrame` that loops through the years and concatenates results. - Save each year's raw data as CSV in `data/raw/acled_ethiopia_{year}.csv`. - Add logging and basic error handling (HTTP errors, unexpected schema). Provide docstrings and type hints."

### 3.3 Command-line entry point or notebook for data download

Create a simple script/notebook to download all years in one go.

#### Atomic Prompt 3.3 – Data download driver

"Create either a small script (e.g., `scripts/download_acled_data.py`) or a Jupyter notebook (`notebooks/01_download_acled.ipynb`) that: - Imports `fetch_acled_range` from `acled_client`. - Uses `config.START_YEAR` and `config.END_YEAR`. - Downloads all ACLED events for Ethiopia and saves a combined CSV at `data/raw/acled_ethiopia_all_years.csv`. Add progress reporting with `tqdm`. Document how to run it in the README."

---

## 4. Geographic Data for Ethiopia

### 4.1 Acquire administrative boundaries

We need **admin boundaries** to map events by region/zone/woreda.

Steps (manual, then codify):

1. Download an Ethiopia administrative boundary dataset (e.g., regions, zones, woredas) as Shapefile or GeoJSON.
2. Save it under `geo/ethiopia_admin_boundaries/`.

#### Atomic Prompt 4.1 – Geo data loading module

"In `src/data_prep.py`, add functions to load Ethiopia administrative boundaries using `geopandas`. Example functions: - `load_admin1_boundaries()` -> `gpd.GeoDataFrame` - `load_admin2_boundaries()` -> `gpd.GeoDataFrame`. Assume the shapefiles/GeoJSONs are stored under `geo/ethiopia_admin_boundaries/` with clear names (we can adjust once files are placed). Include docstrings explaining that these boundaries will be used for mapping pre/post-Abiy conflict patterns."

### 4.2 Decide admin level for main maps

For core blog visuals, we likely want **Admin 1 (regions)**; optionally Admin 2 (zones) for a deeper dive.

#### Atomic Prompt 4.2 – Config for map admin level

"Extend `config.py` with a setting for preferred mapping admin level (e.g., `MAP_ADMIN_LEVEL = 1` for regions). In `data_prep.py`, ensure that the loading functions can be switched or extended to `admin2` if we later decide to map at zone level."

---

## 5. Cleaning & Harmonizing ACLED Data

### 5.1 Standardize columns and types

Key ACLED fields:

- `event_id_cnty`, `event_date`, `year`, `event_type`, `sub_event_type`
- `actor1`, `actor2`, `admin1`, `admin2`, `location`
- `latitude`, `longitude`, `fatalities`

**Atomic Prompt 5.1 – Implement** `clean_acled_data`

"In `src/data_prep.py`, implement a function `clean_acled_data(df: pd.DataFrame) -> pd.DataFrame` for ACLED Ethiopia data. Steps: - Parse `event_date` as `datetime`. - Ensure numeric types for `year` and `fatalities` (coerce errors to `Nan`, then fill with 0 for `fatalities`). - Standardize column names to `snake_case` (e.g., `event_type`, `sub_event_type`, `admin1`, `admin2`). - Drop duplicate rows. - Filter to `country == "Ethiopia"` just in case. Return a clean DataFrame ready for feature engineering. Document assumptions in the docstring."

### 5.2 Save cleaned dataset

**Atomic Prompt 5.2 – Save cleaned CSV & Parquet**

"Add a small script or notebook (e.g., `notebooks/02_clean_acled.ipynb`) that: - Loads `data/raw/acled_ethiopia_all_years.csv`. - Applies `clean_acled_data`. - Saves the result as both `data/interim/acled_ethiopia_clean.csv` and `data/interim/acled_ethiopia_clean.parquet`. Use logging or print statements to show number of rows before/after cleaning."

---

## 6. Feature Engineering – Pre vs Post Abiy & Conflict Metrics

### 6.1 Period indicators (pre vs post)

We need to tag each event as `pre-Abiy` or `post-Abiy` based on `event_date` and `AB_IY_CUTOFF_DATE`.

**Atomic Prompt 6.1 – Add period features**

"In `src/features.py`, implement a function `add_abiy_period_features(df: pd.DataFrame) -> pd.DataFrame` that: - Imports `AB_IY_CUTOFF_DATE` from `config` . -

*Adds a boolean column `is_post_abiy` (True if `event_date` >= `cutoff`). - Adds a categorical column `period` with values "pre\_abiy" and "post\_abiy". - Optionally add `years_since_cutoff` (float) with years relative to cutoff (negative before, positive after). Return an updated DataFrame."*

## 6.2 Aggregate metrics per month and per region

We want time-series and region-level metrics:

- Events per month
- Fatalities per month
- Events/fatalities by `event_type` and `sub_event_type`
- Events per month per `admin1`

### Atomic Prompt 6.2 – Monthly & regional aggregates

*"In `src/features.py`, implement: - `aggregate_monthly(df: pd.DataFrame) -> pd.DataFrame` that groups by `event_month = event_date.to_period('M')` (or year/month columns) and computes total events, total fatalities, plus separate counts for key `event_type` categories. - `aggregate_monthly_by_admin1(df: pd.DataFrame) -> pd.DataFrame` that groups by `admin1` and `month`. Ensure both functions carry a `period` column ('pre\_abiy'/'post\_abiy') and return tidy DataFrames ready for plotting."*

## 6.3 Actor categorization (optional but useful)

Define broad actor categories (e.g., state forces, non-state armed groups, protesters, civilians).

### Atomic Prompt 6.3 – Actor classification skeleton

*"In `src/features.py`, add a helper `classify_actor(actor: str) -> str` that maps ACLED actors into high-level categories: e.g., `state_forces`, `non_state_armed_group`, `protesters`, `civilians`, `other`. Then add a function `add_actor_categories(df: pd.DataFrame) -> pd.DataFrame` that applies this to `actor1` and `actor2` (e.g., `actor1_type`, `actor2_type`). Start with simple keyword rules and keep it easy to extend later."*

---

# 7. Time-Series Analysis (National Level)

## 7.1 Core time-series plots

Key figures (national level):

1. **Monthly events over time (all types)** with a vertical line at 2018-04-02.
2. **Monthly fatalities over time** with the same vertical line.
3. **Events by event\_type** (stacked or faceted) comparing pre/post visually.

### **Atomic Prompt 7.1 – Implement time-series analysis module**

"In `src/analysis_time_series.py`, implement functions that: - Take the monthly aggregate DataFrame and produce summary statistics comparing pre and post periods (mean events/month, mean fatalities/month, etc.). - Return small DataFrames suitable for tables (e.g., `summary_pre_post_overall`, `summary_pre_post_by_event_type`). Include docstrings clearly stating that these functions support the 'Conflict in Ethiopia: Before and After Abiy Ahmed' blog."

## **7.2 Publication-quality plots**

Use `src/visualization.py` for styling.

### **Atomic Prompt 7.2 – Visualization style utilities**

"In `src/visualization.py`, create helper functions to enforce a consistent, publication-quality style using Matplotlib and optionally Seaborn. Requirements: - Set a readable font size and aspect ratio. - Ensure axis labels and titles are clear and informative. - Provide a function `set_publication_style()` that can be called at the start of each plotting function. - Implement a function `add_abiy_cutoff_vline(ax)` that draws a vertical line at the cutoff date on time-series plots and labels it. Use vector-friendly formats (e.g., saving as PNG and SVG/ PDF)."

### **Atomic Prompt 7.3 – Implement time-series plotting functions**

"In `src/visualization.py`, implement plotting functions: - `plot_monthly_events(df_monthly, output_path)` - `plot_monthly_fatalities(df_monthly, output_path)` Each should: - Produce a national time-series plot with `event_month` on the x-axis and counts on the y-axis. - Call `add_abiy_cutoff_vline(ax)` to show the pre/post boundary. - Save the figure to `reports/figures/` with meaningful filenames (e.g., `fig_monthly_events_pre_post_abiy.png`). Ensure the plots are publication-quality (good labels, legend, grid, etc.)."

---

## **8. Spatio-Temporal Analysis**

### **8.1 Link events to admin regions**

We need to spatially join ACLED points to admin boundaries to compute **events per region**.

### **Atomic Prompt 8.1 – Spatial join**

"In `src/analysis_spatiotemporal.py`, implement a function `assign_admin_units(df_events: pd.DataFrame, gdf_admin: gpd.GeoDataFrame) -> gpd.GeoDataFrame` that: - Converts the events DataFrame to a GeoDataFrame using

*latitude/longitude and CRS EPSG:4326. - Performs a spatial join with the admin boundaries to assign each event to an admin1 (and optionally admin2). - Returns a GeoDataFrame of events with added admin attributes. Document assumptions and how to handle events that fall outside boundaries (e.g., drop or mark as unknown)."*

## 8.2 Regional intensity maps (pre vs post)

Construct maps showing **events per admin1** in pre vs post periods.

Key figures:

1. **Choropleth of total events per region (pre-Abiy).**
2. **Choropleth of total events per region (post-Abiy).**
3. **Choropleth of change (post - pre) in events per region**, normalized by years in each period or population (if available).

### Atomic Prompt 8.2 – Aggregations for mapping

"In `src/analysis_spatiotemporal.py`, implement:  
- `aggregate_admin1_period(gdf_events: gpd.GeoDataFrame) -> gpd.GeoDataFrame` that groups by `admin1` and `period` to compute counts of events and fatalities. - Optionally compute per-year rates to account for different lengths of pre/post periods. Return a GeoDataFrame merged with the admin1 geometry, suitable for choropleth mapping."

### Atomic Prompt 8.3 – Choropleth mapping functions

"In `src/visualization.py`, add functions to create Ethiopia choropleth maps: - `plot_admin1_events_pre_post(gdf_admin_agg, output_prefix)` that: - Produces side-by-side maps for `pre_abiy` and `post_abiy` showing events per admin1. - Uses a colorbar, consistent color scale, and clear titles. - Saves maps as high-resolution PNG and SVG/PDF in `reports/figures/`. - Optionally implement `plot_admin1_change_map(gdf_admin_agg, output_path)` to map the change in event counts between periods. Ensure maps are legible and publication-ready (legend, north arrow optional, scale of colors suitable for print)."

## 8.3 Simple spatio-temporal animation (optional)

Optionally, create a **GIF** or **MP4** showing the evolution of monthly events across Ethiopia.

### Atomic Prompt 8.4 – (Optional) animated map

"If feasible, implement an optional function in `analysis_spatiotemporal.py` (or a notebook) that creates a simple spatio-temporal animation: - For each month, map events as points or admin1 counts. - Combine frames into a GIF (using `imageio` or `matplotlib.animation`). - Save the animation to `reports/figures/anim_ethiopia_conflict_timeline.gif`. This is optional but should be coded in a way that can be skipped if dependencies are missing."

---

## 9. Statistical Comparison – Pre vs Post

### 9.1 Descriptive comparisons

We want clear descriptive statistics:

- Average events per month (pre vs post)
- Average fatalities per month (pre vs post)
- Distribution across event types and regions.

#### Atomic Prompt 9.1 – Summary tables

*"In `src/analysis_statistical.py`, implement functions to compute summary statistics: - `summarize_pre_post_overall(df_monthly)` -> `pd.DataFrame` - `summarize_pre_post_by_event_type(df_monthly)` -> `pd.DataFrame` that calculate mean, median, standard deviation of events and fatalities per month in pre vs post periods. Save these as CSV in `reports/tables/` for direct inclusion in the blog (e.g., `tbl_pre_post_overall.csv`, `tbl_pre_post_by_event_type.csv`)."*

### 9.2 Interrupted time series (simple model)

Estimate a simple **interrupted time-series model** at the national level:

- Outcome: events per month.
- Predictors: time trend, post-Abiy dummy, optional interaction.

#### Atomic Prompt 9.2 – ITS regression

*"In `src/analysis_statistical.py`, implement a function `fit_interrupted_time_series(df_monthly)` -> `dict` that: - Constructs a time index (e.g., months since start). - Creates a `post_abiy` dummy. - Fits a simple OLS regression using `statsmodels`: `events_per_month ~ time + post_abiy` (+ `timepost_abiy` if desired). - Returns a dictionary with key coefficients and p-values. - Optionally, export a formatted table (e.g., as LaTeX or markdown) summarizing the model to `reports/tables/its_results.md`. Include comments noting limitations and that this is descriptive, not a causal identification strategy."\**

---

## 10. Central Orchestration Script / Notebook

To make it easy to re-run everything, create an orchestration script or notebook that calls each step in order.

#### Atomic Prompt 10.1 – End-to-end pipeline notebook

*"Create a notebook notebooks/99\_pipeline\_run.ipynb (or a scripts/run\_pipeline.py) that: - Downloads (or loads cached) ACLED Ethiopia data. - Cleans the data. - Adds features (period indicators, aggregates). - Runs time-series, spatio-temporal, and statistical analyses. - Generates all key figures and tables and saves them into reports/figures/ and reports/tables/. Use clear markdown cells to describe each step so a future reader (for EthiopiaConflictAnalytics.org) can understand the pipeline."*

---

## 11. Figure & Table Specification for the Blog

This section tells Cursor **exactly which outputs** to generate for the blog.

### 11.1 Figures

1. **Figure 1 – Monthly conflict events in Ethiopia, from five years before the cutoff date to the most recent available year**
2. Data: df\_monthly national series.
3. Visual: line plot, vertical line at 2018-04-02, shaded bands or annotation for pre/post.
4. **Figure 2 – Monthly conflict fatalities in Ethiopia, from five years before the cutoff date to the most recent available year**
5. Similar to Figure 1, but fatalities.
6. **Figure 3 – Composition of conflict event types, pre vs post Abiy (national)**
7. Bar charts (or stacked bar) comparing share of protests, battles, violence against civilians, etc. pre vs post.
8. **Figure 4 – Regional distribution of conflict events, pre vs post Abiy (Admin 1)**
9. Two choropleth maps side by side.
10. **Figure 5 – Change in conflict intensity by region (post minus pre events per year)**
11. Single choropleth map with diverging color scale.
12. **Optional Figure 6 – Simple ITS model fit**
13. Overlay fitted values from the ITS regression on the monthly events plot.

#### Atomic Prompt 11.1 – Implement figure generation wrapper

*"In src/visualization.py, add a function generate\_all\_blog\_figures(df\_monthly, gdf\_admin\_events, its\_results) that:*

- Calls the individual plotting functions (time-series, composition, maps). - Saves figures using consistent filenames matching the blog numbering (e.g., `fig01_monthly_events.png`, `fig02_monthly_fatalities.png`, etc.). - Returns a dictionary mapping figure labels (e.g., "FIG1") to file paths. This will make it easy to reference the figures when drafting the blog post."

## 11.2 Tables

Core tables to export to `reports/tables/`:

1. **Table 1 – Summary of monthly conflict events and fatalities, pre vs post Abiy**
2. **Table 2 – Event type distribution pre vs post Abiy (share of total events)**
3. **Table 3 – Regional averages of events per month, pre vs post Abiy (Admin 1)**
4. **Table 4 – Interrupted time-series regression coefficients (national events per month)**

### Atomic Prompt 11.2 – Table export utilities

"In `src/analysis_statistical.py` (or `src/visualization.py` if more convenient), implement helper functions to: - Save key summary DataFrames (pre/post overall, by event type, by admin1) as both CSV and markdown tables. - Name files consistently (e.g., `tbl01_pre_post_overall.csv`, `tbl01_pre_post_overall.md`, etc.). Ensure the markdown tables have clean column names and rounded numeric values (2-3 decimal places) so they can be pasted directly into the blog or a static site generator."

---

## 12. Blog Draft Support (Quantitative Sections)

Once the pipeline runs, we want to help draft the quantitative sections of the blog using the outputs.

### 12.1 Narrative stubs for Cursor to generate text

#### Atomic Prompt 12.1 – Generate quantitative narrative

"After the analysis is complete and all figures/tables exist, read the summary tables and, based strictly on those quantitative results, draft a narrative for the blog section titled `Quantitative Trends Before and After Abiy Ahmed`. The narrative should: - Describe the overall change in conflict events and fatalities per month (pre vs post) with approximate percentages. - Highlight which event types grew or declined most. - Identify which regions saw the largest increases or decreases in conflict activity. - Reference figures and tables by their labels (e.g., "Figure 1", "Table 1") but do not invent findings not supported by the data. Keep the tone analytical, neutral, and concise (800-1,000 words)."

### 12.2 Uncertainty and limitations paragraph

#### Atomic Prompt 12.2 – Limitations section

*"Using knowledge of the ACLED dataset and the results produced, draft a short section titled Data Limitations and Caveats for the blog. Cover at least: - Reporting biases and under-reporting. - Differences in period length (pre vs post) and how we adjusted (e.g., per-month or per-year rates). - The descriptive (non-causal) nature of pre/post comparisons. Limit this section to 2-4 paragraphs, in neutral language."*

---

## 13. Quality Control & Reproducibility

### 13.1 Basic tests

#### Atomic Prompt 13.1 – Add lightweight tests

*"Create a tests/ folder with a few simple tests using pytest (or a minimal custom test script) that: - Check that ACLED data loads and has expected columns. - Verify that add\_abiy\_period\_features correctly classifies dates before and after the cutoff. - Ensure monthly aggregation functions return non-empty DataFrames with expected columns. These tests can use small synthetic DataFrames to avoid hitting the API."*

### 13.2 Reproducible run description

#### Atomic Prompt 13.2 – README pipeline section

*"Update the README.md with a How to Reproduce the Analysis section that lists, in order: 1. Create .env with ACLED credentials. 2. Download or place Ethiopia admin boundary data in geo/. 3. Run the data download step. 4. Run cleaning & feature engineering. 5. Run the analysis pipeline (time-series, spatio-temporal, statistics). 6. Inspect figures and tables under reports/. Provide exact commands or notebook names where relevant."*

---

## 14. Optional Extensions (If Time Allows)

These are extras Cursor can implement later, but the core pipeline should work without them.

1. **Population-adjusted rates** using regional population to compute events per 100,000.
2. **Alternative cutoffs** (e.g., first full year of Abiy's government) for robustness.
3. **Sub-period analysis** (e.g., Tigray war, Amhara conflict) nested inside the post-Abiy era.

#### Atomic Prompt 14.1 – Configuration hooks for extensions

*"In config.py, add optional configuration flags (e.g., USE\_POPULATION\_RATES, ALT\_CUTOFF\_DATES) and ensure the analysis modules are written in a way that these can be activated later without major refactoring. You do not need to fully implement the extensions now, but add TODO comments showing where they would plug in."*

---

## 15. Final Checklist for Cursor AI

Before considering the pipeline complete, confirm that:

1. **Data ingestion** from ACLED works and is cached in `data/raw/`.
2. **Cleaning** produces a consistent schema in `data/interim/`.
3. **Features** correctly classify pre/post Abiy and build monthly & regional aggregates.
4. **Time-series analysis** runs and saves figures (events & fatalities) with the cutoff annotated.
5. **Spatio-temporal analysis** generates at least two regional choropleths (pre vs post) and optionally a change map.
6. **Statistical analysis** outputs summary tables and a simple ITS regression table.
7. **All key figures and tables** are saved under `reports/figures/` and `reports/tables/` with clear filenames.
8. **README** explains how to reproduce everything.

### Final Atomic Prompt - Self-check for Cursor AI

*"Review the entire project and confirm that: - The ACLED client, cleaning, feature engineering, time-series, spatio-temporal, and statistical analysis modules are implemented with docstrings and type hints. - Key figures and tables required for the Conflict in Ethiopia: Before and After Abiy Ahmed blog are generated and stored under reports/. - The README provides clear reproduction steps. If any piece is missing or fragile (e.g., hard-coded paths, magic numbers), propose and implement improvements to make the pipeline robust and maintainable."*